

```
In [1]: # Step 1: Data Preparation
```

```
In [2]: import pandas as pd
```

```
In [5]: # Load each dataset (since they are in the same folder, no path needed)
movies = pd.read_csv('movies.csv')
ratings = pd.read_csv('ratings.csv')
tags = pd.read_csv('tags.csv')
links = pd.read_csv('links.csv')
```

```
In [6]: # View first few rows of each
print("Movies:")
display(movies.head())
```

Movies:

	movied	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

```
In [7]: print("\nRatings:")
display(ratings.head())
```

Ratings:

	userId	movied	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931

```
In [8]: print("\nTags:")
display(tags.head())
```

Tags:

	userId	movied	tag	timestamp
0	2	60756	funny	1445714994
1	2	60756	Highly quotable	1445714996
2	2	60756	will ferrell	1445714992
3	2	89774	Boxing story	1445715207
4	2	89774	MMA	1445715200

```
In [9]: print("\nLinks:")
display(links.head())
```

Links:

	movied	imdbId	tmdbdId
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

```
In [10]: # Step 2: Feature Engineering
# Merge ratings with movies
merged_df = pd.merge(ratings, movies, on='movied', how='left')
# Merge with tags
```

```
In [11]: merged_df = pd.merge(merged_df, tags, on=['movieId', 'userId'], how='left')
```

```
# Merge with links (optional)
merged_df = pd.merge(merged_df, links, on='movieId', how='left')
# Display a few rows of the final dataset
merged_df.head()
```

Out[11]:

	userId	movieId	rating	timestamp_x	title	genres	tag	timestamp_y	imdbId	tmdbId
0	1	1	4.0	964982703	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	NaN	NaN	114709	862.0
1	1	3	4.0	964981247	Grumpier Old Men (1995)	Comedy Romance	NaN	NaN	113228	15602.0
2	1	6	4.0	964982224	Heat (1995)	Action Crime Thriller	NaN	NaN	113277	949.0
3	1	47	5.0	964983815	Seven (a.k.a. Se7en) (1995)	Mystery Thriller	NaN	NaN	114369	807.0
4	1	50	5.0	964982931	Usual Suspects, The (1995)	Crime Mystery Thriller	NaN	NaN	114814	629.0

In [12]:

```
# STEP 5: DATA CLEANING
# Check for missing values
print("Missing values per column:\n", merged_df.isnull().sum())
# Drop duplicates if any
merged_df.drop_duplicates(inplace=True)
# Convert timestamp columns to datetime format
merged_df['timestamp_x'] = pd.to_datetime(merged_df['timestamp_x'], unit='s', errors='coerce')
merged_df['timestamp_y'] = pd.to_datetime(merged_df['timestamp_y'], unit='s', errors='coerce')
# Rename for clarity
merged_df.rename(columns={'timestamp_x': 'rating_timestamp', 'timestamp_y': 'tag_timestamp'}, inplace=True)
print("\n✓ Data cleaned and timestamps converted successfully.")
merged_df.head()
```

Missing values per column:

```
userId          0
movieId         0
rating          0
timestamp_x     0
title           0
genres          0
tag             99201
timestamp_y     99201
imdbId          0
tmdbId          13
dtype: int64
```

✓ Data cleaned and timestamps converted successfully.

Out[12]:

	userId	movieId	rating	rating_timestamp	title	genres	tag	tag_timestamp	imdbId	t
0	1	1	4.0	2000-07-30 18:45:03	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	NaN	NaN	114709	
1	1	3	4.0	2000-07-30 18:20:47	Grumpier Old Men (1995)	Comedy Romance	NaN	NaN	113228	15602.0
2	1	6	4.0	2000-07-30 18:37:04	Heat (1995)	Action Crime Thriller	NaN	NaN	113277	
3	1	47	5.0	2000-07-30 19:03:35	Seven (a.k.a. Se7en) (1995)	Mystery Thriller	NaN	NaN	114369	
4	1	50	5.0	2000-07-30 18:48:51	Usual Suspects, The (1995)	Crime Mystery Thriller	NaN	NaN	114814	

In [13]:

```
# STEP 6: FEATURE ENGINEERING
```

In [14]:

```
import numpy as np
```

In [15]:

```
# [1] Extract release year from the movie title
merged_df['release_year'] = merged_df['title'].str.extract(r'\((\d{4})\)').astype(float)
```

In [18]:

```
# [2] Count number of genres per movie
merged_df['genre_count'] = merged_df['genres'].apply(lambda x: len(x.split('|'))) if isinstance(x, str) else 0
```

```
In [19]: # [3] Calculate how old the movie is (as of 2025)
merged_df['movie_age'] = 2025 - merged_df['release_year']

In [20]: # [4] Average rating per movie
movie_avg_rating = merged_df.groupby('movieId')['rating'].mean().reset_index(name='avg_rating')
merged_df = merged_df.merge(movie_avg_rating, on='movieId', how='left')

In [21]: # [5] Number of ratings per movie
movie_rating_count = merged_df.groupby('movieId')['rating'].count().reset_index(name='rating_count')
merged_df = merged_df.merge(movie_rating_count, on='movieId', how='left')

In [22]: # [6] Rating hour (to see what time of day people tend to rate)
merged_df['rating_hour'] = merged_df['rating_timestamp'].dt.hour

In [24]: # [7] Weekday rating (0 = Monday)
merged_df['rating_weekday'] = merged_df['rating_timestamp'].dt.day_name()

In [28]: print("✓ Feature engineering complete! Here are the new columns:\n")
print(merged_df[['title', 'release_year', 'genre_count', 'movie_age', 'avg_rating', 'rating_count', 'rating_hour']])

✓ Feature engineering complete! Here are the new columns:

          title  release_year  genre_count  movie_age \
0      Toy Story (1995)      1995.0         5     30.0
1  Grumpier Old Men (1995)      1995.0         2     30.0
2            Heat (1995)      1995.0         3     30.0
3  Seven (a.k.a. Se7en) (1995)      1995.0         2     30.0
4   Usual Suspects, The (1995)      1995.0         3     30.0
...           ...
102672             Split (2017)      2017.0         3      8.0
102673  John Wick: Chapter Two (2017)      2017.0         3      8.0
102674            Get Out (2017)      2017.0         1      8.0
102675            Logan (2017)      2017.0         2      8.0
102676  The Fate of the Furious (2017)      2017.0         4      8.0

      avg_rating  rating_count  rating_hour
0       3.920930        215          18
1       3.245283         53          18
2       3.946078        102          18
3       3.980392        204          19
4       4.252404        208          18
...           ...
102672     3.333333         6          21
102673     4.307692        13          22
102674     3.633333         15          19
102675     4.241379        29          21
102676     2.333333         3          21

[102677 rows x 7 columns]


In [29]: # STEP 7: EXPLORATORY DATA ANALYSIS (EDA)

In [30]: import matplotlib.pyplot as plt
import seaborn as sns

In [31]: sns.set_style('whitegrid')
plt.rcParams['figure.figsize'] = (10,5)

In [32]: # [1] Distribution of Ratings
sns.histplot(merged_df['rating'], bins=10, kde=True)
plt.title("Distribution of Movie Ratings")
plt.xlabel("Rating")
plt.ylabel("Count")
plt.show()
```


Distribution of Movie Ratings



```
In [33]: # [2] Top 10 Movies by Number of Ratings
top_movies = merged_df.groupby('title')['rating'].count().sort_values(ascending=False).head(10)
sns.barplot(x=top_movies.values, y=top_movies.index, palette='coolwarm')
plt.title("Top 10 Most Rated Movies")
plt.xlabel("Number of Ratings")
plt.ylabel("Movie Title")
plt.show()
```

C:\Users\MR. ROWLINGS\AppData\Local\Temp\ipykernel_7996\3726818738.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.




```
In [34]: # [3] Average Rating per Genre
genre_avg = merged_df.groupby('genres')['rating'].mean().sort_values(ascending=False).head(10)
sns.barplot(x=genre_avg.values, y=genre_avg.index, palette='mako')
plt.title("Average Rating by Genre")
plt.xlabel("Average Rating")
plt.ylabel("Genre")
plt.show()
```


C:\Users\MR. ROWLINGS\AppData\Local\Temp\ipykernel_7996\3992641893.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=genre_avg.values, y=genre_avg.index, palette='mako')
```



```
In [36]: # [4] Relationship between Number of Ratings and Average Rating
sns.scatterplot(data=merged_df.drop_duplicates('movieId'),
x='rating_count', y='avg_rating')
plt.title("Rating Count vs. Average Rating per Movie")
plt.xlabel("Number of Ratings")
plt.ylabel("Average Rating")
plt.show()
```




```
In [38]: # [5] Most Common Tags
top_tags = merged_df['tag'].value_counts().head(10)
sns.barplot(x=top_tags.values, y=top_tags.index, palette='viridis')
plt.title("Top 10 Most Common Tags")
plt.xlabel("Count")
plt.ylabel("Tag")
plt.show()
```


C:\Users\MR. ROWLINGS\AppData\Local\Temp\ipykernel_7996\447396862.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_tags.values, y=top_tags.index, palette='viridis')
```



```
In [39]: # [6] Ratings by Day of the Week
sns.countplot(x='rating_weekday', data=merged_df,
order=['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'])
plt.title("Number of Ratings by Day of Week")
plt.xticks(rotation=45)
plt.show()
```



Key Insights

1. Ratings Distribution: Most users give ratings between 3.0 and 4.5, showing a positive bias.
2. Most Rated Movies: A few popular titles dominate the ratings count.
3. Genres: Certain genres like Drama and Action tend to have higher average ratings.
4. Tags: User tags show engagement trends (e.g., "classic", "funny", "sci-fi" appear often).
5. Activity Pattern: Ratings peak around weekends.
6. Correlation: Movies with more ratings often have slightly lower average ratings — a popularity trade-off. These insights help in designing a recommendation system by identifying popular genres, user preferences, and engagement patterns.

```
In [40]: merged_df.to_csv("cleaned_movielen.csv", index=False)
print("✓ Cleaned dataset saved as 'cleaned_movielen.csv'")
```

✓ Cleaned dataset saved as 'cleaned_movielen.csv'

```
In [ ]:
```